

☆ Image Style Transfer ☆

Gatys et. al. 2016

1. Introduction

Texture analysis → Efros & Freeman (correspondence map)
→ Hertzman et. al. (image analogies)
→ Ashikhmin (transferring the high-frequency texture info + preserve coarse scale)
→ Lee et. al. improved + edge orientation info.

remarkable results but limitation = only low-level image features of the target img. to inform the texture transfer

◆ a style transfer algo. should be able to extract the semantic img. content from the target img. (e.g. objects & scenery.) and then inform a texture transfer procedure to render the semantic content of the target img. in the style of the source img.

➡ to separate content from style, extract high-level semantic info. by using Convolutional Neural Networks (CNNs).

introduced: A Neural Algo. of Artistic Style.

✚ combines a parametric texture model based on CNNs with a method to invert their img. ~~represent~~ representations.

2. Deep Image Representations

• basis → VGG network which is for object recognition & localization.

VGG-19 } 16 convolutional & 5 pooling layers
total 19 layers = 16 conv + 3 fully connected



● network is normalized by scaling the weights
→ the mean activation of each conv. filter over imgs. and positions is equal to one.

★ re-scaling can be done for the VGG net. without changing its output; because it contains only rectifying linear activation (ReLU) funcs. and no normalization or pooling over feature maps.

- not used any of fully connected layers.
- for img. synthesis avg. pooling is slightly better than max. pooling.

2.1 Content Representation

- each layer in network defines a non-linear filter bank
- + complexity of layer increases with the position of layer.
- + hence, input img. \vec{x} is encoded in each layer of CNN by filter responses.

N_l → distinct layers

M_l → size of each feature map (height x width)

F^l → respons in layer l


↓
matrix

$$F^l \in \mathbb{R}^{N_l \times M_l}$$

, F_{ij}^l → activation of the i^{th} filter at position j in layer l .



perform gradient descent on a white noise img. to find another img. that matches the features responses of original img.



let \vec{p} & \vec{x} be original img., P^l & F^l their respective feature representations in layer l . ^{white-noise}

Then, squared-error loss:

$$\mathcal{L}_{\text{content}}(\vec{p}, \vec{x}, l) = \frac{1}{2} \sum_{i,j} (F_{ij}^l - P_{ij}^l)^2$$

Derivative of loss func. wrt. activations in layer l :

$$\frac{\partial \mathcal{L}_{\text{content}}}{\partial F_{ij}^l} = \begin{cases} (F_{ij}^l - P_{ij}^l) & \text{if } F_{ij}^l > 0 \\ 0 & \text{if } F_{ij}^l < 0 \end{cases}$$

\vec{x} can be computed using standard error back-propagation

low layers reproduce exact pixel values of original img.
so, higher layers are preferred \rightarrow as content representation.

2.2. Style Representation

- for style, a feature space to capture texture info.
- + this feature map can be built on top of the filter responses in any layer of net.
- + consists of correlations between different filter responses \Rightarrow represented by Gram matrix.

$G^l \in \mathbb{R}^{N_l \times M_l}$, $G_{ij}^l \rightarrow$ inner product between vectorised feature map i and j in l .

$$G_{ij}^l = \sum_k F_{ij}^l \cdot F_{jk}^l$$



- feature correlations of multiple layers \rightarrow stationary, + texture info, not global averaged (multiscale representation)

◆ visualize info. by style feature spaces:

* * * * * \rightarrow by using gradient descent from a white noise img. to minimise the mean-squared distance between entries of Gram matrices from original img. & Gram matrices of img. to be generated.

let \vec{a} & \vec{x} be original img. and generated img.

A^l & G respectively style representation in layer l .

Then, contribution of layer l to total loss:

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{ij} (G_{ij}^l - A_{ij}^l)$$

Total style loss:

$$L_{\text{style}}(\vec{a}, \vec{x}) = \sum_{l=0}^L \overset{\text{weighting factor}}{w_l} \cdot E_l$$

Derivative of E_l :

$$\frac{\partial E_l}{\partial F_{ij}^l} = \begin{cases} \frac{1}{N_l^2 M_l^2} ((F^l)^T \cdot (G^l - A^l))_{ij} & \text{if } F_{ij}^l > 0 \\ 0 & \text{if } F_{ij}^l < 0 \end{cases}$$

2.3. Style Transfer

The loss function that is minimized:

$$L_{\text{total}}(\vec{p}, \vec{a}, \vec{x}) = \alpha \cdot L_{\text{content}}(\vec{p}, \vec{x}) + \beta \cdot L_{\text{style}}(\vec{a}, \vec{x})$$




- α & β are weighting factors.
- $\frac{\partial h_{total}}{\partial \vec{x}}$ can be used as input for some numerical optimisation strategy.
- here L-BFGS is used (best for img. synthesis)
- always resized style imgs to same size as the content.
- not regularised synthesis results with img. priors.

3. Results

3.1. Trade-off between content & style matching

3.2. Effect of different layers of the CNN

- matching the style representations up to higher layers in the network preserves local img. structures on increasingly large scale.  'conv1_1', 'conv2_1', 'conv3_1', 'conv4_1' and 'conv5_1'

- lower layers \rightarrow more content from original img.

3.3. Initialisation of Gradient Descent

- white noise is used but one can also initialize it with content or style img; but it causes bias.

3.4. Photorealistic Style Transfer



4. Discussion

- most limiting factor is resolution of the synthesised img.
- dimensionality of the optimisation problem } grow
+ the number of units in the CNN } linearly
so, speed depends on resolution with # of pixels.

512x512 pixels } in this paper → prohibits online & interactive applications
Nvidia K40 GPU }
one hour

- separation of img. content from style is not well-defined problem.

Bonus: L-BFGS : Limited-memory BFGS

• an optimization algorithm that approximates BFGS^{*} algorithm (Broyden-Fletcher-Goldfarb-Shanno) using a limited amount of computer memory. family: quasi-Newton methods.

④ an iterative method for solving unconstrained ~~minimization~~ non-linear optimization. $O(n^2)$

- popular algo. for parameter estimation in ML.
- target problem is to minimize $f(x)$ over constrained values.

Quasi-Newton Method: used to either find local maxima & minima of functions or zeroes. They can be used if the Jacobian or Hessian is unavailable or too expensive.

- L-BFGS aka "the algo. of choice" for fitting log-linear (MaxEnt) and conditional random fields with ℓ_2 -regularization.



◆ a way of finding a (local) minimum of an objective func., making use of objective func. values and the gradient of the obj. func.